

# Training Large Language Models

## How to reduce memory and compute

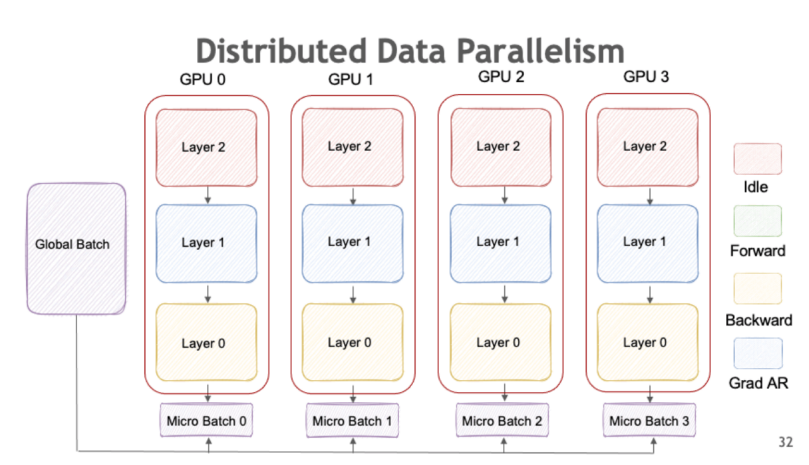
### Learning goals

- Learn about different techniques to reduce compute and memory
- Learn about distributed training with data/tensor parallelism
- Learn about FlashAttention

# DISTRIBUTED TRAINING

- Training LLMs faster on many GPUs
- Avoiding OOM issues
- **Data parallelism:** split the data on different model replicas
- **Tensor parallelism:** split model parameters accross GPUs

# DATA PARALLELISM (1) (ANIMATED GIF!)



Source: Nvidia

# DATA PARALLELISM (2)

- **Data Splitting:**

- The dataset is divided into smaller chunks, and each chunk is assigned to a different processing unit (e.g., GPU or CPU) on different nodes
- Each node processes a different subset of the data in parallel, reducing the overall training time

- **Model Replication:**

- Each processing unit has a replica of the neural network model
- These replicas are trained independently on their respective data subsets

- **Gradient Aggregation:**

- After each forward and backward pass, gradients are computed locally on each node


# DATA PARALLELISM (2)

- The gradients are then averaged (or summed) across all nodes to ensure that each model replica receives the same gradient update
- **Parameter Synchronization:**
  - The model parameters (weights and biases) are updated synchronously across all nodes
  - This ensures that all model replicas remain consistent with each other after each update step

# TENSOR PARALLELISM (1) (ANIMATED GIF!)

## Tensor Parallelism (Intra-Layer)


MLP



Split weights  $W$  and  $V^T$  across multiple GPUs

$$W = [W_1, W_2, W_3, W_4] \quad V = \begin{bmatrix} V_1 \\ V_2 \\ V_3 \\ V_4 \end{bmatrix}$$

Tensor Parallel = 4 MLP



Source: Nvidia

# TENSOR PARALLELISM (2)

- **Model Partitioning:**

- The model's layers or tensors are split across multiple devices
- Different parts of the model are assigned to different devices, enabling them to work on separate portions of the computations simultaneously

- **Forward and Backward Passes:**

- During the forward pass, each device processes its portion of the tensors with intermediate results passed between devices
- In the backward pass, gradients are computed in the reverse order, again with necessary data transfers between devices

- **Parameter Updates:**

- Parameter updates can be performed independently on each device for the parameters they own
- After each update step, the devices synchronize to ensure consistency across the distributed model


# FlashAttention

## Fast and Memory-Efficient Exact Attention with IO-Awareness

- Fast
  - 15 % faster than BERT
  - 3x faster than GPT-2
  - 2.4x faster than Megatron-LM
- Memory-efficient
  - Reducing from  $O(n^2)$  to  $O(n)$
- Exact
  - Same as “vanilla attention”, not an approximation
- IO aware
  - Reducing memory load/store operations



# GPU MEMORY HIERARCHY




Source: Dao et al. (2022)

# COMPUTING CONSIDERATIONS

- GPU compute has been growing faster than memory bandwidth
  - GPU has to wait for data
- Transformer operations are memory-bound
  - Elementwise operations with high memory access
- IO aware means reducing memory load/store operations
- FlashAttention implements the following:
  - Operation fusion to reduce memory access
  - Tiling or chunking the softmax matrix into blocks
  - Recomputation for better memory utilization

# OPERATION FUSION



Source: [https://horace.io/brrr\\_intro.html](https://horace.io/brrr_intro.html)

# LIMITATIONS AND PROSPECTS

- FlashAttention requires writing attention to CUDA language
  - A new CUDA kernel for each new attention implementation
  - CUDA is lower-level than PyTorch
  - Implementation may not be transferable accross GPUs
- Towards IO-Aware Deep Learning
  - Extending beyonde attention
- Multi-GPU IO-Aware Methods
  - FlashAttention computation may be parallelizable accross multiple GPUs